

#Program 1

# Plot a boxplot for `price` vs `cut` from the dataset `diamond.csv`. Which of the categories under `cut` have the highest median price?

```
import pandas as pd
import matplotlib.pyplot as plt


# Read the dataset
df = pd.read_csv("diamonds.csv")

print(df.head())

# Create the boxplot
plt.figure(figsize=(8, 6))
df.boxplot(column='price', by='cut')
plt.xlabel('Cut')
plt.ylabel('Price')
plt.title('Price vs Cut')
plt.suptitle('') # Remove the default title
plt.show()
```

```
0  1  0.23  Ideal  E  SI2  61.5  55.0  326  3.95  3.98
1  2  0.21  Premium  E  SI1  59.8  61.0  326  3.89  3.84
2  3  0.23  Good  E  VS1  56.9  65.0  327  4.05  4.07
3  4  0.29  Premium  I  VS2  62.4  58.0  334  4.20  4.23
4  5  0.31  Good  J  SI2  63.3  58.0  335  4.34  4.35
```

```
z
0  2.43
1  2.31
2  2.31
3  2.63
4  2.75
<Figure size 800x600 with 0 Axes>
```



#Program 2

```
'''Create a frequency table (one-way table) for the variable cut from the dataset diamond.csv. What is the frequency for the cut type Ideal'''
```

```
import pandas as pd

# Read the dataset
df = pd.read_csv("diamonds.csv")

# Create the frequency table
frequency_table = df['cut'].value_counts()

# Print the frequency table
print(frequency_table)

# Access the frequency for the cut type "Ideal"
ideal_frequency = frequency_table['Ideal']
```

```
print("Frequency for cut type 'Ideal':", ideal_frequency)
```

```

Ideal      21551
Premium    13791
Very Good  12082
Good       4906
Fair       1610
Name: cut, dtype: int64
Frequency for cut type 'Ideal': 21551

```

```
#Program 3
```

```
# Show the subplot of the diamond carat weight distribution.
```

```
import pandas as pd
import matplotlib.pyplot as plt
```

```
# Read the dataset
df = pd.read_csv("diamonds.csv")
```

```
print(df['carat'])
```

```
# Create a subplot with 1 row and 2 columns
fig, axs = plt.subplots(1, 2, figsize=(12, 6))
```

```
# Plot the histogram of carat weight on the first subplot
axs[0].hist(df['carat'], bins=30, edgecolor='black')
axs[0].set_xlabel('Carat Weight')
axs[0].set_ylabel('Frequency')
axs[0].set_title('Carat Weight Distribution')
```

```
# Plot the boxplot of carat weight on the second subplot
axs[1].boxplot(df['carat'], vert=False)
axs[1].set_yticklabels('')
axs[1].set_xlabel('Carat Weight')
axs[1].set_title('Carat Weight Distribution')
```

```
# Adjust the spacing between subplots
plt.subplots_adjust(wspace=0.3)
```

```
# Show the plot
plt.show()
```

```
0      0.23
1      0.21
2      0.23
```

```
#Program 4
```

```
# Show the subplot of diamond depth distribution.
```

```
import pandas as pd
import matplotlib.pyplot as plt
```

```
# Read the dataset
df = pd.read_csv("diamonds.csv")
print(df['depth'])
```

```
# Create a subplot with 1 row and 2 columns
fig, axs = plt.subplots(1, 2, figsize=(12, 6))
```

```
# Plot the histogram of diamond depth on the first subplot
axs[0].hist(df['depth'], bins=30, edgecolor='black')
axs[0].set_xlabel('Depth')
axs[0].set_ylabel('Frequency')
axs[0].set_title('Diamond Depth Distribution')
```



```
# Plot the boxplot of diamond depth on the second subplot
axs[1].boxplot(df['depth'], vert=False)
axs[1].set_yticklabels('')
axs[1].set_xlabel('Depth')
axs[1].set_title('Diamond Depth Distribution')
```

```
# Adjust the spacing between subplots
plt.subplots_adjust(wspace=0.4)
```

```
# Show the plot
plt.show()
```

```
0      61.5
1      59.8
2      56.9
3      62.4
4      63.3
...
53935   60.8
53936   63.1
53937   62.8
53938   61.0
53939   62.2
```

```
Name: depth, Length: 53940, dtype: float64
```



```
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import r2_score
```

```

# Read the dataset
df = pd.read_csv("diamonds.csv")

# Define the input features (X) and target variable (y)
X = df[['carat', 'cut', 'clarity', 'depth']]
y = df['price']

# Convert categorical variables to numerical using one-hot encoding
X_encoded = pd.get_dummies(X)

print("X_encoded=", X_encoded)
# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X_encoded, y, test_size=0.2, random_state=42)

# Build the linear regression model
model = LinearRegression()

model.fit(X_train, y_train)

# Make predictions on the testing set
y_pred = model.predict(X_test)

# Calculate the accuracy (R-squared score)
accuracy = r2_score(y_test, y_pred)
print("Accuracy:", accuracy)

```

X_encoded=	carat	depth	cut_Fair	cut_Good	cut_Ideal	cut_Premium	\
0	0.23	61.5	0	0	1	0	
1	0.21	59.8	0	0	0	1	
2	0.23	56.9	0	1	0	0	
3	0.29	62.4	0	0	0	1	
4	0.31	63.3	0	1	0	0	
...	...	...	...	...	...	...	
53935	0.72	60.8	0	0	1	0	
53936	0.72	63.1	0	1	0	0	
53937	0.70	62.8	0	0	0	0	
53938	0.86	61.0	0	0	0	1	
53939	0.75	62.2	0	0	1	0	

	cut_Very Good	clarity_I1	clarity_IF	clarity_SI1	clarity_SI2	\
0	0	0	0	0	1	
1	0	0	0	1	0	
2	0	0	0	0	0	
3	0	0	0	0	0	
4	0	0	0	0	1	
...	...	...	...	...	...	
53935	0	0	0	1	0	
53936	0	0	0	1	0	
53937	1	0	0	1	0	
53938	0	0	0	0	1	
53939	0	0	0	0	1	

	clarity_VS1	clarity_VS2	clarity_VS1	clarity_VS2
0	0	0	0	0
1	0	0	0	0
2	1	0	0	0
3	0	1	0	0
4	0	0	0	0
...	...	...	...	...
53935	0	0	0	0
53936	0	0	0	0
53937	0	0	0	0
53938	0	0	0	0
53939	0	0	0	0

```

[53940 rows x 15 columns]
LinearRegression()
Accuracy: 0.8966651780887069

```

